

Incomplete factorization

Karl Meerbergen

KU Leuven

October 16, 2013

Literature

- James Demmel, Applied numerical Linear Algebra, SIAM, 1997.
Gives a good overview of preconditioning techniques.
- Anne Greenbaum, Iterative methods for solving linear systems, SIAM, 1997.
Gives an excellent overview of Krylov methods. The book also briefly discusses preconditioning techniques.
- Youcef Saad, Iterative methods for sparse linear systems, SIAM, 2003.
Gives an excellent overview of Krylov methods. The book also discusses incomplete factorization techniques.

LU-factorization

- Linear algebra:

$$A = LU$$

with L unit lower triangular and U upper triangular

- Numerical linear algebra:
 - ▶ Pivoting for stability.
 - ▶ Fill-in when matrices are sparse.
 - ▶ High computational cost $O(n^3)$ for a **full** matrix.
- Solving linear system $Ax = b$ with backtransformation:
 - ▶ $Ly = b$
 - ▶ $Ux = y$

LU-factorization

- Algorithm:

```
for  $k = 1, \dots, n$  do  
  for  $i = k + 1, \dots, n$  do  
     $a_{ik} = a_{ik} / a_{kk}$   
    for  $j = k + 1, \dots, n$  do  
       $a_{ij} = a_{ij} - a_{ik} a_{kj}$   
    end for  
  end for  
end for
```

- There are six (6) different forms to write this algorithm: the six permutations of 3 loops.
- This algorithm is the *kij* form
- Different forms are used in different circumstances

Right-looking LU-factorization

- kji form is column oriented and right-looking.

- Algorithm:

```
for  $k = 1, \dots, n$  do  
  for  $j = k + 1, \dots, n$  do  
     $a_{jk} = a_{jk} / a_{kk}$   
    for  $i = k + 1, \dots, n$  do  
       $a_{ij} = a_{ij} - a_{ik} a_{kj}$   
    end for  
  end for  
end for
```

Other forms

- The six forms:

type	orientation	updating
kji	column	right-looking
jki	column	left-looking
jik	column	left-looking
kij	row	right-looking
ikj	row	
ijk	row	

- Naming convention from the FORTRAN world where matrices are column oriented.

Sparse matrix storage formats

- Two formats are in use (and a few variations):
 - ▶ Coordinate format
 - ▶ Compressed sparse column/row format
- Illustrate for the following matrix

$$\begin{bmatrix} -3.5 & 1.4 & & 8 & -2 \\ & -3 & & & \\ & & 3 & 2 & \\ & & & 1 & -5 \\ & & & & -5 \end{bmatrix}.$$

Sparse matrix storage formats

$$\begin{bmatrix} -3.5 & 1.4 & & 8 & -2 \\ & -3 & & & \\ & & 3 & 2 & \\ & & & 1 & -5 \\ & & & & -5 \end{bmatrix}.$$

- Coordinate format

- ▶ rows: the row number of the non-zero
- ▶ columns: the column number of the non-zero
- ▶ values: the value of the non-zero

A possible COO representation for the example is

rows:	1	2	3	3	1	1	4	1	4	5
columns:	1	2	3	4	2	4	4	5	5	5
values:	-3.5	-3	3	2	1.4	8	1	-2	-5	-5

The non-zeroes can be stored in any order: this is one of the great advantages of the format.

Sparse matrix storage formats

$$\begin{bmatrix} -3.5 & 1.4 & & 8 & -2 \\ & -3 & & & \\ & & 3 & 2 & \\ & & & 1 & -5 \\ & & & & -5 \end{bmatrix}.$$

- Compressed sparse column (row) format Sorting the coordinate format by column we obtain the following data:

rows r:	1	1	2	3	3	4	1	1	4	5
columns c:	1	2	2	3	4	4	4	5	5	5
values v:	-3.5	1.4	-3	3	2	1	8	-2	-5	-5

Replace the columns array by an array that points to the beginning of each column in the array of row numbers and values:

columns c:	1	2	4	5	8	11
------------	---	---	---	---	---	----

Sparse matrix storage formats

- j -th element of `column` is a pointer to the beginning and end of a column
- Fast access of a column is thus possible
- It simplifies and speeds up matrix operations such as matrix vector product.
- For matrix assembly, the coordinate format is much more convenient since elements can be added in any order.
- For matrix assembly, the compressed format requires the use of insertion in an array, which is expensive.

LU factorization of a sparse matrix

- $A = LU$: $L + U$ is usually denser than A .

$$\begin{bmatrix} \times & & \times & \times & & \times \\ & \times & & \times & & \\ \times & & \times & & \times & \\ \times & \times & & \times & & \times \\ & & \times & & \times & \times \\ \times & & & \times & \times & \times \end{bmatrix} A = LU \Rightarrow \begin{bmatrix} \times & & \times & \times & & \times \\ & \times & & \times & & \\ \times & & \times & \circ & \times & \circ \\ \times & \times & \circ & \times & \circ & \times \\ & & \times & \circ & \times & \times \\ \times & & \circ & \times & \times & \times \end{bmatrix}$$

- Factorization: $A = (L + D)D^{-1}(D + U)$ with L strictly lower diagonal and U strictly upper triangular
- Factorization:

for $k = 1, \dots, n$ **do**

for all (i, j) for which $a_{i,k}a_{k,j} \neq 0$ **do**

$$f_{i,j} = -a_{i,k}a_{k,j}/a_{k,k}$$

$$a_{ij} = a_{i,j} + f_{i,j}$$

end for

end for

How does it work?

- Factorization: $A = (L + D)D^{-1}(D + U)$
- i, j element:

$$a_{i,j} = \sum_{k=1}^{\min(i,j)} l_{i,k} d_k^{-1} u_{k,j}$$

$$l_{i,j} = a_{i,j} - \sum_{k=1}^{j-1} l_{i,k} d_k^{-1} u_{k,j} \quad (u_{j,j} d_j^{-1} = 1)$$

$$u_{i,j} = a_{i,j} - \sum_{k=1}^{i-1} l_{i,k} d_k^{-1} u_{k,j} \quad (l_{i,i} d_i^{-1} = 1)$$

or when A stores L , D^{-1} and U :

$$a_{i,j} = a_{i,j} - \sum_{k=1}^{\min(i,j)} a_{i,k} a_{k,j} / a_{k,k}$$

How does it work?

- Algorithm:

for $k = 1, \dots, n$ **do**

for all (i, j) for which $a_{i,k}a_{k,j} \neq 0$ **do**

$$f_{i,j} = -a_{i,k}a_{k,j}/a_{k,k}$$

$$a_{ij} = a_{ij} + f_{i,j}$$

end for

end for

- Factorize row k and column k and update all other elements.

<i>Done</i>		$a_{1,k}$	<i>Done</i>
		\vdots	
$a_{k,1}$	\cdots	$a_{k,k}$	
<i>Done</i>			<i>Update</i>

Variations

- Some prefer the following decomposition:

$$A = (L + D)D^{-1}(U + D)$$

with L strictly lower triangular and U strictly upper triangular

- Classical factorization in numerical analysis text books:

$$A = (L + I)(U + D)$$

- LDU factorization:

$$A = (L + I)D(U + I)$$

- Cholesky factorization:

$$A = (L + D)(L + D)^T$$

where A is symmetric positive definite.

Incomplete factorization

- Let S be the sparsity pattern of A : $S = \{(i, j) : a_{i,j} \neq 0\}$.

- Algorithm:

```
for  $k = 1, \dots, n$  do  
  for all  $(i, j) \in S$  for which  $a_{i,k} a_{k,j} \neq 0$  do  
     $f_{i,j} = -a_{i,k} a_{k,j} / a_{k,k}$   
     $a_{ij} = a_{i,j} + f_{i,j}$   
  end for  
end for
```

M-matrix

- LU-factorization exists for a positive definite matrix (= Cholesky factorization, no pivoting required)
- Incomplete LU factorization may fail, even when the matrix is non-singular
- Sufficient condition for success:

Definition ($A \in \mathbb{C}^{n \times n}$ is an *M-matrix*)

- 1 $a_{ii} > 0$, for $i = 1, \dots, n$,
 - 2 $a_{ij} \leq 0$, $i, j = 1, \dots, n$, $i \neq j$, and
 - 3 A is nonsingular and all elements of A^{-1} are non-negative.
- When A is a matrix with non-positive off-diagonal entries, this is equivalent with any of the following statements:
 - 1 The eigenvalues of A have positive real parts.
 - 2 $A + A^T$ is positive definite.
 - 3 All principal minors of A are *M-matrices*.

Example

- Example:
- Laplacian:

$$L = \left[\begin{array}{ccc|ccc} 4 & -1 & & -1 & & \\ -1 & 4 & -1 & & -1 & \\ & -1 & 4 & & & -1 \\ \hline -1 & & & 4 & -1 & \\ & -1 & & -1 & 4 & -1 \\ & & -1 & & -1 & 4 \end{array} \right]$$

- Eigenvalues are 1.5858, 3.0000, 3.5858, 4.4142, 5.0000, and 6.4142
- Inverse

$$L^{-1} = \frac{1}{2415} \begin{bmatrix} 712 & 225 & 68 & 208 & 120 & 47 \\ 225 & 780 & 225 & 120 & 255 & 120 \\ 68 & 225 & 712 & 47 & 120 & 208 \\ 208 & 120 & 47 & 712 & 225 & 68 \\ 120 & 255 & 120 & 225 & 780 & 225 \\ 47 & 120 & 208 & 68 & 225 & 712 \end{bmatrix}$$

Example

- Matrix

$$A = \begin{bmatrix} 3 & & -1 & -1 & & -1 \\ & 2 & & -1 & & \\ -1 & & 3 & & -1 & \\ -1 & -1 & & 2 & & -1 \\ & & -1 & & 3 & -1 \\ -1 & & & -1 & -1 & 4 \end{bmatrix}$$

has eigenvalues 0.2246, 1.5249, 2.5493, 3.2019, 4.1355, 5.3638 and has non-positive off-diagonal elements, so it is an M matrix.

- (Full) Cholesky factorization:

$$\tilde{L} = L + D = \begin{bmatrix} 1.7321 & & & & & \\ 0 & 1.4142 & & & & \\ -0.5774 & 0 & 1.6330 & & & \\ -0.5774 & -0.7071 & -0.2041 & 1.0607 & & \\ 0 & 0 & -0.6124 & -0.1179 & 1.6159 & \\ -0.5774 & 0 & -0.2041 & -1.2964 & -0.7908 & 1.1485 \end{bmatrix}$$

Example

- Matrix

$$A = \begin{bmatrix} 3 & & -1 & -1 & & -1 \\ & 2 & & -1 & & \\ -1 & & 3 & & -1 & \\ -1 & -1 & & 2 & & -1 \\ & & -1 & & 3 & -1 \\ -1 & & & -1 & -1 & 4 \end{bmatrix}$$

has eigenvalues 0.2246, 1.5249, 2.5493, 3.2019, 4.1355, 5.3638 and has non-positive off-diagonal elements, so it is an M matrix.

- incomplete Cholesky factorization:

$$\tilde{L} = \begin{bmatrix} 1.7321 & & & & & \\ 0 & 1.4142 & & & & \\ -0.5774 & 0 & 1.6330 & & & \\ -0.5774 & -0.7071 & 0 & 1.0801 & & \\ 0 & 0 & -0.6124 & 0 & 1.6202 & \\ -0.5774 & 0 & 0 & -1.2344 & -0.6172 & 1.3274 \end{bmatrix}$$

Example

- incomplete Cholesky factorization:

$$\tilde{L} = \begin{bmatrix} 1.7321 & & & & & \\ 0 & 1.4142 & & & & \\ -0.5774 & 0 & 1.6330 & & & \\ -0.5774 & -0.7071 & 0 & 1.0801 & & \\ 0 & 0 & -0.6124 & 0 & 1.6202 & \\ -0.5774 & 0 & 0 & -1.2344 & -0.6172 & 1.3274 \end{bmatrix}$$

with

$$R = A - \tilde{L}\tilde{L}^T = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & -0.3333 & 0 & -0.3333 \\ 0 & 0 & -0.3333 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & -0.3333 & 0 & 0 & 0 \end{bmatrix}.$$

The eigenvalues of $\tilde{L}^{-1}A\tilde{L}^{-T}$ are 1.0000, 0.5305, 1.3536, 1.0000, 1.0000, 1.0000. The spectrum is pretty well clustered around one.

Theorem

- ① *For any M-matrix A and for any sparsity pattern S that contains the main diagonal elements, there exist unique L and U (L lower triangular with ones on the main diagonal and U upper triangular) where $L + U$ have the sparsity pattern S , so that*

$$A - (L + D)D^{-1}(U + D) = R$$

where $r_{ij} = 0$ for all $(i, j) \in S$.

- ② *For any symmetric M matrix and (symmetric) sparsity pattern S containing the main diagonal, there exists a unique lower triangular matrix L so that*

$$A - (L + D)D^{-1}(U + D) = R$$

where $r_{ij} = 0$ for all $(i, j) \in S$.

Modified ILU (MILU)

- Modify D so that $Ae = (L + D)D^{-1}(D + U)e$
- Motivation: preconditioner has the same 'mass' as the original matrix.
- Mathematical explanation: for second order self-adjoint elliptic equations, the ILU preconditioner produces $\text{cond}((LU)^{-1}A) = O(h^{-2})$ while MILU produces $\text{cond}((LU)^{-1}A) = O(h^{-1})$
- Exact factorization of element $a_{i,j}$:

$$\sum_{j=1}^n a_{i,j} = \sum_{j=1}^n \sum_{k=1}^{\min(i,j)} l_{i,k} d_k^{-1} u_{k,j}$$

- With dropping strategy:

$$\sum_{j=1, (i,j) \in S}^n \sum_{k=1}^{\min(i,j)} l_{i,k} d_k^{-1} u_{k,j} + \sum_{j=1, (i,j) \notin S}^n \sum_{k=1}^{\min(i,j)} l_{i,k} d_k^{-1} u_{k,j}$$

- Accumulate lost elements in D (the elements that belong to $R = A - L \cdot U$)

Modified ILU (MILU)

```
1: for  $k = 1, \dots, n$  do
2:   for  $(i, j)$  for which  $a_{i,k}a_{k,j} \neq 0$  do
3:     Compute  $f_{i,j} = -a_{i,k}a_{k,j}/a_{k,k}$ .
4:     if  $(i, j) \in S$  then
5:        $a_{i,j} = a_{i,j} + f_{i,j}$ 
6:     else
7:        $a_{i,i} = a_{i,i} + f_{i,j}$ 
8:     end if
9:   end for
10: end for
```

Other alternatives than MILU

- Relaxed ILU (RILU)

```
1: for  $k = 1, \dots, n$  do
2:   for  $(i, j)$  for which  $a_{i,k}a_{k,j} \neq 0$  do
3:     Compute  $f_{i,j} = -a_{i,k}a_{k,j}/a_{k,k}$ .
4:     if  $(i, j) \in S$  then
5:        $a_{i,j} = a_{i,j} + f_{i,j}$ 
6:     else
7:        $a_{i,i} = a_{i,i} + \alpha f_{i,j}$  with  $0 < \alpha < 1$ 
8:     end if
9:   end for
10: end for
```


- Helmholtz equation:

$$-\nabla^2 u - k^2 u = f \quad \Rightarrow \quad (A - k^2 I) \mathbf{u} = \mathbf{f}$$


$A - k^2 I$ is indefinite and usually not an M matrix. Instead, build the preconditioner for $A + k^2 I$

Helmholtz equation


- Spectrum of $A - k^2 I$



- Spectrum of $(A + k^2 I)^{-1}(A - k^2 I)$



- Spectrum of $(A - k^2 I + i2k^2 I)^{-1}(A - k^2 I)$



Additional fill-in

- (M)ILU use the sparsity pattern of A for S
- The number of nonzero terms in the residual $R = A - (L + D)D^{-1}(U + D)$ can be reduced by using a larger set S .
- ILU does not work well for matrices that are not M -matrices.
- Various ideas:
 - ▶ Level of fill
 - ▶ Thresholding small absolute values
 - ▶ Pivoting
 - ▶ Other orderings

ILU with level of fill

- Fill-in = zero elements of A that become nonzero during factorization.
- There is an intuitive feeling that fill-in resulting from other fill-in elements is less important than fill-in coming from non-zero elements of A .

Definition

Level of fill

$$\text{Levfill}(a_{i,j}) = \min(\text{Levfill}(a_{i,j}), \max(\text{Levfill}(a_{i,k}), \text{Levfill}(a_{k,j})) + 1)$$

- Example

$$\begin{bmatrix} 0 & & 0 & 0 & & \\ & 0 & & 0 & & \\ 0 & & 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 & 2 & 0 \\ & 0 & & 2 & 0 & 0 \\ 0 & 1 & & 0 & 0 & 0 \end{bmatrix}$$

ILU with level of fill

- Compute LevFill:

- 1: $\text{LevFill}(a_{i,j}) = 0$ for all $(i,j) \in S$ and $\text{LevFill}(a_{i,j}) = \infty$ for all $(i,j) \notin S$.
- 2: **for** $k = 1, \dots, n$ **do**
- 3: **for** all (i,j) **do**
- 4: $\text{Levfill}(a_{i,j}) =$
 $\min(\text{Levfill}(a_{i,j}), \max(\text{Levfill}(a_{i,k}), \text{Levfill}(a_{k,j})) + 1)$
- 5: **end for**
- 6: **end for**

Example

- Matrix

$$A = \begin{bmatrix} 3 & & -1 & -1 & & -1 \\ & 2 & & -1 & & \\ -1 & & 3 & & -1 & \\ -1 & -1 & & 2 & & -1 \\ & & -1 & & 3 & -1 \\ -1 & & & -1 & -1 & 4 \end{bmatrix}$$

- ILU(1):

$$\tilde{L} = \begin{bmatrix} 1.7321 & & & & & \\ 0 & 1.4142 & & & & \\ -0.5774 & 0 & 1.6330 & & & \\ -0.5774 & -0.7071 & -0.2041 & 1.0607 & & \\ 0 & 0 & -0.6124 & 0 & 1.6202 & \\ -0.5774 & 0 & -0.2041 & -1.2964 & -0.6944 & 1.2093 \end{bmatrix}$$

The eigenvalues of $(\tilde{L}\tilde{L}^T)^{-1}A$ are 1.0000, 0.8323, 1.0000, 1.0782, 1.0000, and 1.0000.

Example

- Note that


$$A - \tilde{L}\tilde{L}^T = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & -0.1250 & 0 \\ 0 & 0 & 0 & -0.1250 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

- For ILU(0), we had

$$R = A - \tilde{L}\tilde{L}^T = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & -0.3333 & 0 & -0.3333 \\ 0 & 0 & -0.3333 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & -0.3333 & 0 & 0 & 0 \end{bmatrix}.$$

Example

- Model for acoustic radiation to infinity



- ▶ 72,976 unknowns
- ▶ Finite elements
- ▶ Infinite elements for acoustic radiation
- ▶ Load consists of 'rotating modes'.
- ▶ $Ax = f \quad A(\omega) = K + i\omega C - \omega^2 M$

- Timings for different preconditioners

ILU(0)		ILU(1)	
P_1	P_2	P_1	P_2
72.03	50.01	89.03	88.28

with $P_1 = \text{ILU for } A$ and $P_2 = \text{ILU for } A((1 - i)\omega)$.

- ILU(1) requires 1.8 times the memory of ILU(0).

Thresholding and pivoting

- $\text{ILU}(p)$ works well for M -matrices. Many problems are not M -matrices
- Throw away small elements in the matrix and use a direct solver: this does not work well when all elements are of the same order
- Throw away small elements in the factorization (ILUT)

1: **for** $k = 1, \dots, n$ **do**

2: Compute the elements in row k and column k .

3: Throw away all elements smaller than $\tau \|a_{k,:}\|$.

4: Keep the largest p elements.

5: **end for**

- ILUTP: ILUT with pivoting

- ▶ Pivoting can help significantly in improving the quality of the preconditioner. As for direct methods, it is advantageous to pivot elements with large modulus to the main diagonal.

Multilevel techniques

- Also see domain decomposition
- ILU is hard to parallelize: the algorithm is sequential in principle.
- This parallelization of ILU is a problem similar to the parallelization of sparse direct methods. This is discussed in another lecture: see the choice of renumbering and the resulting elimination tree.
- For ILU, we can use an approach similar to domain decomposition:

$$\begin{bmatrix} A_1 & & & A_{1,4} \\ & A_2 & & A_{2,4} \\ & & A_3 & A_{3,4} \\ A_{4,1} & A_{4,2} & A_{4,3} & A_4 \end{bmatrix} \approx \begin{bmatrix} I & & & \\ & I & & \\ & & I & \\ \tilde{A}_{4,1} & \tilde{A}_{4,2} & \tilde{A}_{4,3} & I \end{bmatrix} \cdot \begin{bmatrix} A_1 & & & A_{1,4} \\ & A_2 & & A_{2,4} \\ & & A_3 & A_{3,4} \\ & & & S \end{bmatrix}$$

- We have that $A_{4,i}A_i^{-1} \approx \tilde{A}_{4,i}$ and

$$S \approx A_4 - \sum \tilde{A}_{4,i}A_i^{-1}A_{i,4}$$

Multilevel techniques

- In AMRS and IMF, a multilevel approach is adopted, where S is repeatedly treated in the same way.


$$\left[\begin{array}{c|cc} D_1 & & U_1 \\ \hline L_1 & D_2 & U_2 \\ & L_2 & S_2 \\ & \underbrace{\hspace{1cm}} & \\ & S_1 & \end{array} \right] \approx \left[\begin{array}{c|cc} D_1 & & U_1 \\ \hline \tilde{L}_1 & \tilde{D}_2 & \tilde{U}_2 \\ & \tilde{L}_2 & \tilde{S}_2 \\ & \underbrace{\hspace{1cm}} & \\ & \tilde{S}_1 & \end{array} \right]$$

with $\tilde{L}_1 \approx L_1 D_1^{-1}$, etc.


- ARMS
 - ▶ D , L , U , S are sparse matrices. \tilde{L}_i is a sparse approximation of $A_{i,i+1} D_i^{-1}$.
- IMF
 - ▶ D_i is a block diagonal matrix
 - ▶ L_i is stored as a factored matrix $A_{i,i+1} D_i^{-1}$ where S_i^{-1} is stored explicitly
 - ▶ The only approximation is in the Schur complement.

Multilevel techniques


- Sparse arithmetic is less efficient than dense arithmetic.
- ARMS uses sparse matrix operations
- In IMF, D_i is stored as a dense block diagonal matrix



(a) Throughput of computational primitives.



(b) Number of operations of a BiCgSTAB iteration per second.



(c) Number of BiCgSTAB iterations per second.